

Deep Learning for Audio

Lecture 9

Pavel Severilov

AI Masters

2024

Outline

1. Speaker verification and identification
2. Metric learning
3. Triplet loss
4. Angular softmax
5. ArcFace

Outline

1. Speaker verification and identification
2. Metric learning
3. Triplet loss
4. Angular softmax
5. ArcFace

Metric learning

- ▶ Task: determine how similar are two objects
- ▶ Data:
 - ▶ Supervised: labeled objects
 - ▶ Unsupervised: set of similar or dissimilar pairs
- ▶ Applications:
 - ▶ Few-shot learning
 - ▶ Biometrics
 - ▶ Face recognition (who is on the photo?)
 - ▶ Speaker verification (who is speaking?)
 - ▶ Large amount of rare classes

Speaker recognition: base pipelines

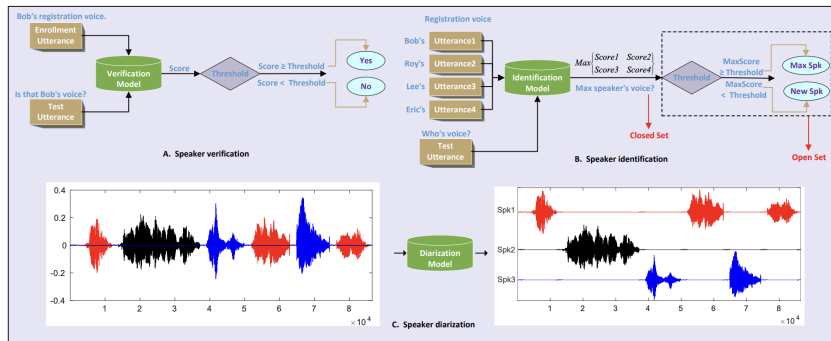



Figure: Flowcharts of speaker verification, speaker identification, and speaker diarization

Speaker verification: metrics



- ▶ $\text{model}(\text{audio1.wav}, \text{audio2.wav}) = 1.03$ - need thresholds
- ▶ False positive - true: different, prediction: same
- ▶ False negative - true: same, prediction: different
- ▶ Equal Error Rate (EER) - the point where $\text{TPR} == \text{FPR}$ for a given model

VoxCeleb2 dataset

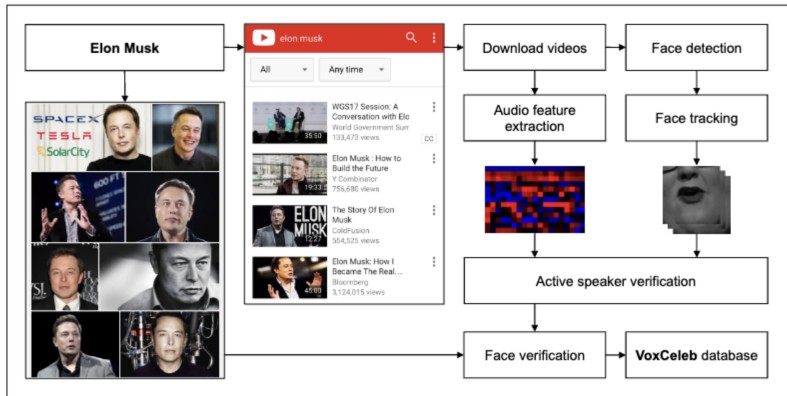



Figure: Data processing pipeline

Outline

1. Speaker verification and identification
2. Metric learning
3. Triplet loss
4. Angular softmax
5. ArcFace

Metric learning idea

a) Original data space




b) Euclidean metric


$d(x, y)$

A diagram showing two points, x (an orange triangle) and y (a black square), with a double-headed arrow between them labeled $d(x, y)$.
$$d(x, y) = \sqrt{\sum_{i=1}^k (x_i - y_i)^2}$$

c) Purpose of metric learning




d) Deep metric learning example*



*Siamese Network

e) Transformed data space



Cosine loss

Idea: “attract” everything that have to be close, “spread” everything that shouldn't be close. Problems?

$$\text{loss}(x, y) = \begin{cases} 1 - \cos(x_1, x_2), & \text{if } y = 1 \\ \max(0, \cos(x_1, x_2) - \text{margin}), & \text{if } y = -1 \end{cases}$$

Contrastive loss

Why should we “spread” everything that shouldn’t be close? Let’s spread only the objects that are too close.

Similarity fn



$$L_+(x_1, x_2) = \frac{1}{4}(1 - E_W)^2$$

$$L_-(x_1, x_2) = \begin{cases} E_W^2 & \text{if } E_W < m \\ 0 & \text{otherwise} \end{cases}$$

Outline

1. Speaker verification and identification
2. Metric learning
3. Triplet loss
4. Angular softmax
5. ArcFace

Triplet loss

Why should we even move anything, if everything is fine?


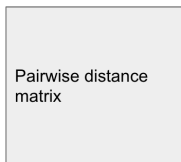



Figure: Easy negative, hard negative, semi-hard negative

$$\mathcal{L}_{\text{triplet}}(\mathbf{x}, \mathbf{x}^+, \mathbf{x}^-) = \sum_{\mathbf{x} \in \mathcal{X}} \max \left(0, \|f(\mathbf{x}) - f(\mathbf{x}^+)\|_2^2 - \|f(\mathbf{x}) - f(\mathbf{x}^-)\|_2^2 + \epsilon \right)$$

Triplet mining: how to sample triplets?

- ▶ Offline (compute all the embeddings on the training set, and then only select hard or semi-hard triplets)
- ▶ Online:



- Don't mine positives, just take each positive with each positive
- For each positive pair (Anchor, Positive) choose negative example so that , loss > 0

$$\|f(x_i^a) - f(x_i^p)\|_2^2 < \|f(x_i^a) - f(x_i^n)\|_2^2$$

Outline

1. Speaker verification and identification
2. Metric learning
3. Triplet loss
4. Angular softmax
5. ArcFace

Angular softmax

Softmax function:

$$p_1 = \frac{e^{\mathbf{W}_1^T \mathbf{x} + b_1}}{e^{\mathbf{W}_1^T \mathbf{x} + b_1} + e^{\mathbf{W}_2^T \mathbf{x} + b_2}}$$

$$p_2 = \frac{e^{\mathbf{W}_2^T \mathbf{x} + b_2}}{e^{\mathbf{W}_1^T \mathbf{x} + b_1} + e^{\mathbf{W}_2^T \mathbf{x} + b_2}}$$

- ▶ Let's remove biases, so that $b_i = 0$. Then, the decision boundary:

$$(\mathbf{W}_1^T - \mathbf{W}_2^T) \mathbf{x} = 0 \rightarrow (||\mathbf{W}_1|| \cos(\theta_1) - ||\mathbf{W}_2|| \cos(\theta_2)) ||\mathbf{x}|| = 0$$


- ▶ Let's normalize weights, so that $||\mathbf{W}_i|| = 1$

$$(||\mathbf{W}_1|| \cos(\theta_1) - ||\mathbf{W}_2|| \cos(\theta_2)) ||\mathbf{x}|| = 0 \rightarrow \cos(\theta_1) - \cos(\theta_2) = 0$$

Angular softmax

$$\cos(\theta_1) - \cos(\theta_2) = 0$$

Example: decision boundary for $W_1 = \left(\frac{\sqrt{2}}{2}, \frac{\sqrt{2}}{2}\right)$, $W_2 = \left(\frac{-1}{2}, \frac{\sqrt{3}}{2}\right)$



Angular softmax

Let's introduce a more restrictive condition (m - positive integer)

- ▶ class 1, if $\cos(m\theta_1) > \cos(\theta_2)$
- ▶ class 2, if $\cos(m\theta_2) > \cos(\theta_1)$



Then the loss function:

$$L = -\log \frac{e^{\|\mathbf{x}^{(n)}\| \cos(m\theta_{y_n}^{(n)})}}{e^{\|\mathbf{x}^{(n)}\| \cos(m\theta_{y_n}^{(n)})} + \sum_{j \neq y_n} e^{\|\mathbf{x}^{(n)}\| \cos(\theta_j^{(n)})}}$$

Angular softmax: meaning of m

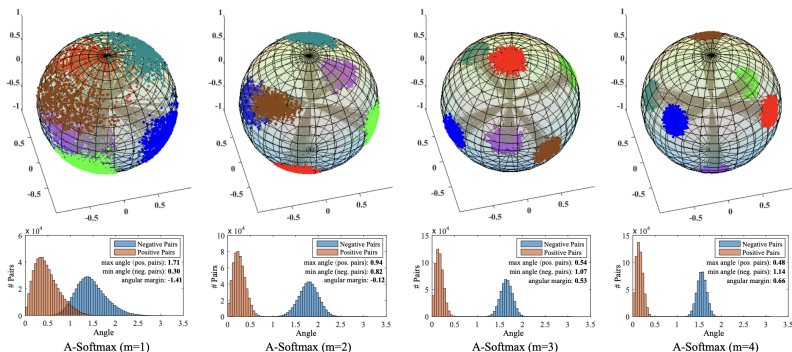


Figure: Visualization of features learned with different m .

- First row: 3D features projected on the unit sphere.
- Second row: the angle distribution of both positive pairs and negative pairs

Outline

1. Speaker verification and identification
2. Metric learning
3. Triplet loss
4. Angular softmax
5. ArcFace

ArcFace

Based on Angular softmax:

- ▶ Remove biases
- ▶ Normalize weights
- ▶ + Normalize features x
- ▶ + Additive m instead of multiplication


Angular softmax loss:

$$L = -\log \frac{e^{\|\mathbf{x}^{(n)}\| \cos(m\theta_{y_n}^{(n)})}}{e^{\|\mathbf{x}^{(n)}\| \cos(m\theta_{y_n}^{(n)})} + \sum_{j \neq y_n} e^{\|\mathbf{x}^{(n)}\| \cos(\theta_j^{(n)})}}$$

Arcface loss:

$$L = -\log \frac{e^{s \cos(\theta_{y_i} + m)}}{e^{s \cos(\theta_{y_i} + m)} + \sum_{j=1, j \neq y_i}^N e^{s \cos \theta_j}}$$

ArcFace: pipeline



ArcFace: geometric interpretation




Figure: Decision margins of different loss functions under binary classification case. The dashed line represents the decision boundary, and the grey areas are the decision margins.